

Архитектура операционной системы

Files, filesystems, vfs

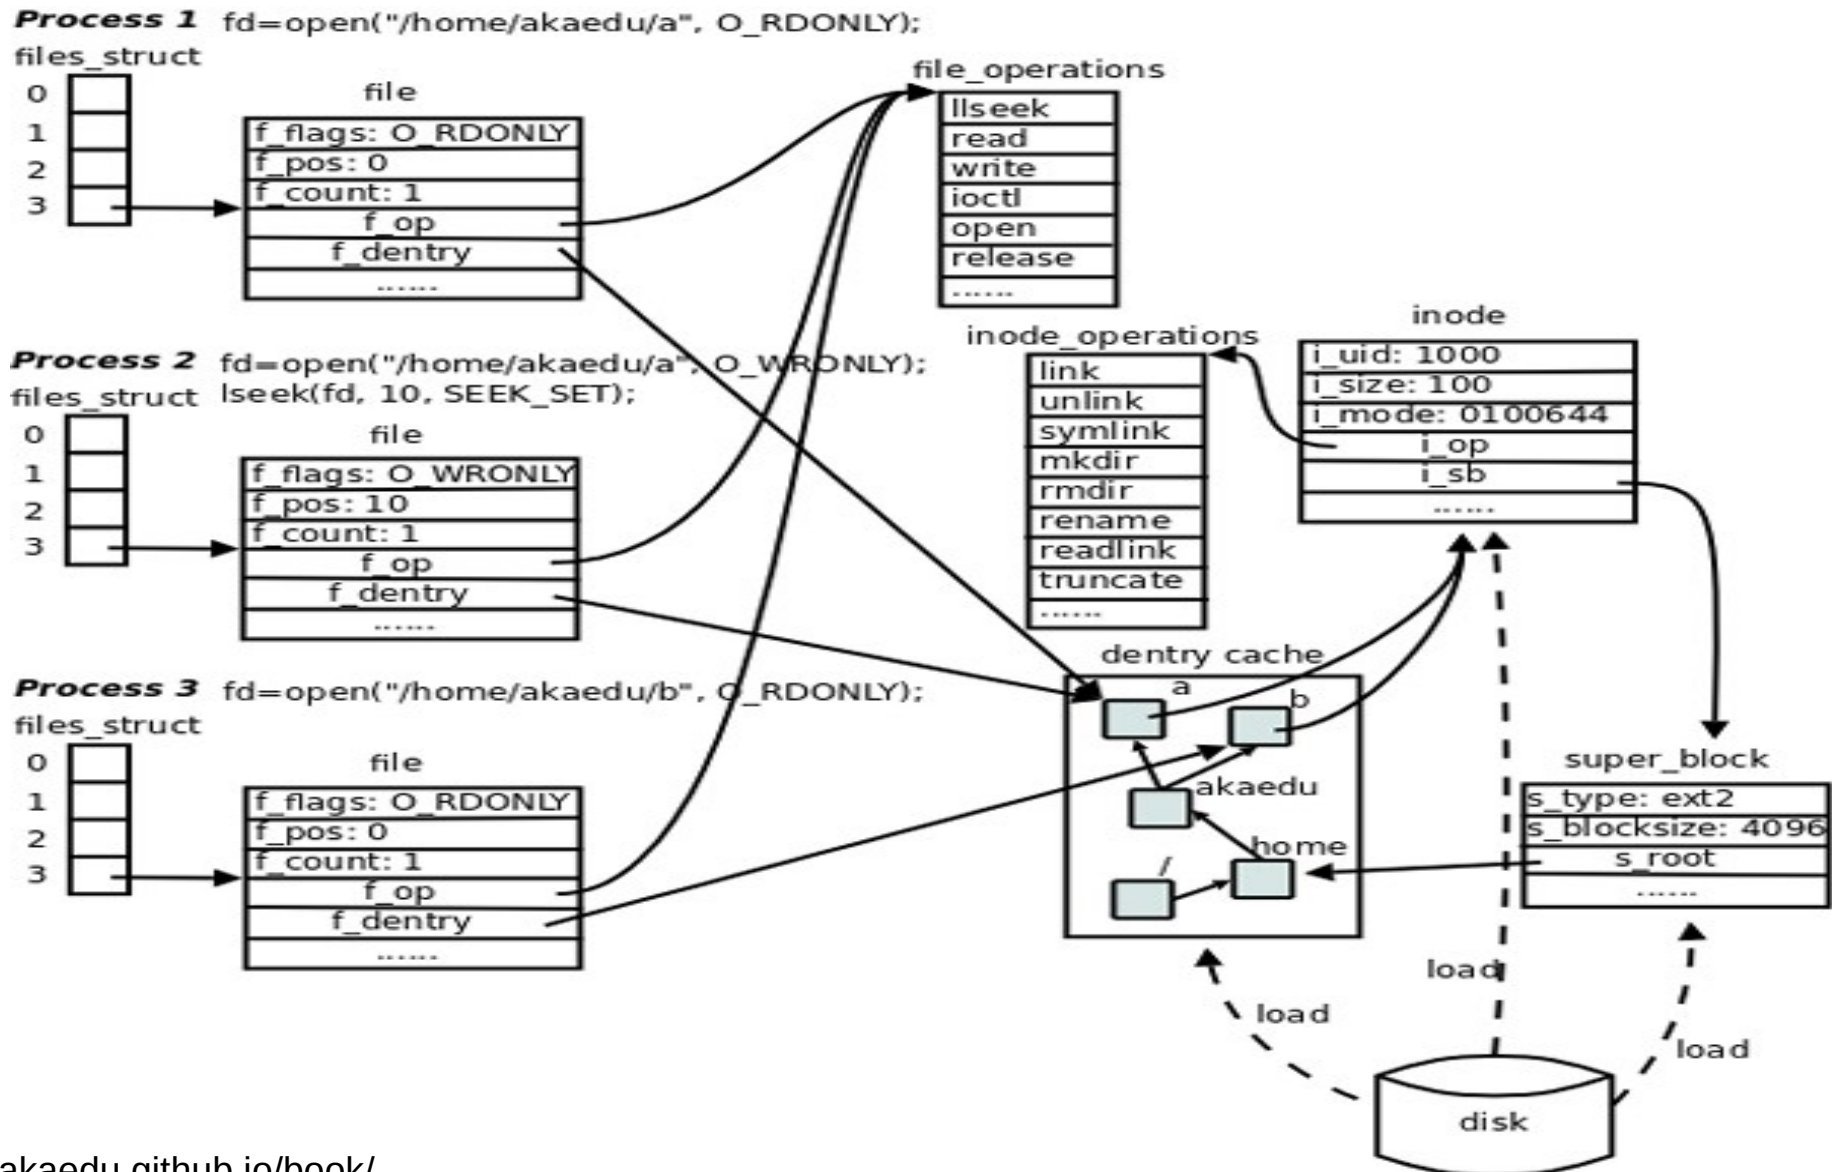
Назначение

- Реализация единого интерфейса для построения файловых систем
- Уровень абстракции от конкретной файловой системы
- Стандартизация доступа к ФС на разных носителях

Структуры данных VFS

- `superblock` – главный (первый) блок файловой системы, описывающий ее параметры (метаинформация о ФС)
- `inode` – индексный узел (метаинформация о файле)
- `dentry` – элемент пути (каталог)
- `file` – объект, представляющий собой открытый файл, связанный с процессом

Связь структур данных



<http://akaedu.github.io/book/>

file_systems → * file_system_type

```
struct file_system_type {
    const char *name;
    int fs_flags;
#define FS_REQUIRES_DEV          1
#define FS_BINARY_MOUNTDATA     2
#define FS_HAS_SUBTYPE          4
#define FS_USERNS_MOUNT         8      /* Can be mounted by usersns root */
#define FS_RENAME_DOES_D_MOVE   32768 /* FS will handle d_move() during rename() internally. */
    struct dentry *(*mount) (struct file_system_type *, int,
                            const char *, void *);
    void (*kill_sb) (struct super_block *);
    struct module *owner;
    struct file_system_type * next;
    struct hlist_head fs_supers;

    struct lock_class_key s_lock_key;
    struct lock_class_key s_umount_key;
    struct lock_class_key s_vfs_rename_key;
    struct lock_class_key s_writers_key[SB_FREEZE_LEVELS];

    struct lock_class_key i_lock_key;
    struct lock_class_key i_mutex_key;
    struct lock_class_key i_mutex_dir_key;
};
```

};

Структуры данных VFS

- `superblock` – главный (первый) блок файловой системы, описывающий ее параметры (метаинформация о ФС)
- `inode` – индексный узел (метаинформация о файле)
- `dentry` – элемент пути (каталог)
- `file` – объект, представляющий собой открытый файл, связанный с процессом

current → namespace → * vfmount

```
struct vfmount {  
    struct dentry *mnt_root;           /* root of the mounted tree */  
    struct super_block *mnt_sb;       /* pointer to superblock */  
    int mnt_flags;  
} __randomize_layout;  
  
struct file; /* forward dec */
```

Current → namespace → * mnt_sb

```

1308
1309 struct super_block {
1310     struct list_head      s_list;           /* Keep this first */
1311     dev_t                 s_dev;           /* search index; _not_ kdev
1312     unsigned char         s_blocksize_bits;
1313     unsigned long         s_blocksize;
1314     loff_t                s_maxbytes;      /* Max file size */
1315     struct file_system_type *s_type;
1316     const struct super_operations *s_op;
1317     const struct dquot_operations *dq_op;
1318     const struct quotactl_ops *s_qcop;
1319     const struct export_operations *s_export_op;
1320     unsigned long         s_flags;
1321     unsigned long         s_iflags;        /* internal SB_I_* flags */
1322     unsigned long         s_magic;
1323     struct dentry          *s_root;
1324     struct rw_semaphore    s_umount;
1325     int                    s_count;
1326     atomic_t              s_active;
1327     #ifdef CONFIG_SECURITY
1328         void                *s_security;
1329     #endif
1330     const struct xattr_handler **s_xattr;

```




```
1770
1777 struct super_operations {
1778     struct inode *(*alloc_inode)(struct super_block *sb);
1779     void (*destroy_inode)(struct inode *);
1780
1781     void (*dirty_inode) (struct inode *, int flags);
1782     int (*write_inode) (struct inode *, struct writeback_control *wbc);
1783     int (*drop_inode) (struct inode *);
1784     void (*evict_inode) (struct inode *);
1785     void (*put_super) (struct super_block *);
1786     int (*sync_fs)(struct super_block *sb, int wait);
1787     int (*freeze_super) (struct super_block *);
1788     int (*freeze_fs) (struct super_block *);
1789     int (*thaw_super) (struct super_block *);
1790     int (*unfreeze_fs) (struct super_block *);
1791     int (*statfs) (struct dentry *, struct kstatfs *);
1792     int (*remount_fs) (struct super_block *, int *, char *);
1793     void (*umount_begin) (struct super_block *);
1794
1795     int (*show_options)(struct seq_file *, struct dentry *);
1796     int (*show_devname)(struct seq_file *, struct dentry *);
1797     int (*show_path)(struct seq_file *, struct dentry *);
1798     int (*show_stats)(struct seq_file *, struct dentry *);
1799 #ifdef CONFIG_QUOTA
```





```
192 struct export_operations {
193     int (*encode_fh)(struct inode *inode, __u32 *fh, int *max_len,
194                     struct inode *parent);
195     struct dentry * (*fh_to_dentry)(struct super_block *sb, struct fid *fid,
196                                     int fh_len, int fh_type);
197     struct dentry * (*fh_to_parent)(struct super_block *sb, struct fid *fid,
198                                     int fh_len, int fh_type);
199     int (*get_name)(struct dentry *parent, char *name,
200                    struct dentry *child);
201     struct dentry * (*get_parent)(struct dentry *child);
202     int (*commit_metadata)(struct inode *inode);
203
204     int (*get_uuid)(struct super_block *sb, u8 *buf, u32 *len, u64 *offset);
205     int (*map_blocks)(struct inode *inode, loff_t offset,
206                      u64 len, struct iomap *iomap,
207                      bool write, u32 *device_generation);
208     int (*commit_blocks)(struct inode *inode, struct iomap *iomaps,
209                          int nr_iomaps, struct iattr *iattr);
210 };
211
212 extern int exportfs_encode_inode_fh(struct inode *inode, struct fid *fid,
213                                    int *max_len, struct inode *parent);
214 extern int exportfs_encode_fh(struct dentry *dentry, struct fid *fid,
```





```
847
848 struct file {
849     union {
850         struct llist_node    fu_llist;
851         struct rcu_head      fu_rcuhead;
852     } f_u;
853     struct path              f_path;
854     struct inode             *f_inode;    /* cached value */
855     const struct file_operations *f_op;
856
857     /*
858      * Protects f_ep_links, f_flags.
859      * Must not be taken from IRQ context.
860      */
861     spinlock_t               f_lock;
862     enum rw_hint              f_write_hint;
863     atomic_long_t             f_count;
864     unsigned int              f_flags;
865     fmode_t                   f_mode;
866     struct mutex              f_pos_lock;
867     loff_t                    f_pos;
868     struct fown_struct        f_owner;
869     const struct cred         *f_cred;
870     struct file_ra_state      f_ra;
```





```
561  /*
562   * Keep mostly read-only and often accessed (especially for
563   * the RCU path lookup and 'stat' data) fields at the beginning
564   * of the 'struct inode'
565   */
566  struct inode {
567      umode_t          i_mode;
568      unsigned short   i_opflags;
569      kuid_t           i_uid;
570      kgid_t           i_gid;
571      unsigned int      i_flags;
572
573  #ifdef CONFIG_FS_POSIX_ACL
574      struct posix_acl  *i_acl;
575      struct posix_acl  *i_default_acl;
576  #endif
577
578      const struct inode_operations *i_op;
579      struct super_block *i_sb;
580      struct address_space *i_mapping;
581
582  #ifdef CONFIG_SECURITY
583      void *i_security;
584  #endif
645      const struct file_operations *i_fop; /* former ->i_op->default_file_ops
```




```
1664 struct file_operations {
1665     struct module *owner;
1666     loff_t (*llseek) (struct file *, loff_t, int);
1667     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
1668     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
1669     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
1670     ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
1671     int (*iterate) (struct file *, struct dir_context *);
1672     int (*iterate_shared) (struct file *, struct dir_context *);
1673     unsigned int (*poll) (struct file *, struct poll_table_struct *);
1674     long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
1675     long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
1676     int (*mmap) (struct file *, struct vm_area_struct *);
1677     int (*open) (struct inode *, struct file *);
1678     int (*flush) (struct file *, fl_owner_t id);
1679     int (*release) (struct inode *, struct file *);
1680     int (*fsync) (struct file *, loff_t, loff_t, int datasync);
1681     int (*fasync) (int, struct file *, int);
1682     int (*lock) (struct file *, int, struct file_lock *);
1683     ssize_t (*sendpage) (struct file *, struct page *, int, size_t, loff_t);
1684     unsigned long (*get_unmapped_area)(struct file *, unsigned long, unsigned
1685     int (*check_flags)(int);
1686     int (*flock) (struct file *, int, struct file_lock *);
```





1703

1704

1705

1706

1707

1708

1709

1710

1711

1712

1713

1714

1715

1716

1717

1718

1719

1720

1721

1722

1723

1724

1725

1726

```
struct inode_operations {
```

```
    struct dentry * (*lookup) (struct inode *, struct dentry *, unsigned int
```

```
    const char * (*get_link) (struct dentry *, struct inode *, struct delay
```

```
    int (*permission) (struct inode *, int);
```

```
    struct posix_acl * (*get_acl)(struct inode *, int);
```

```
    int (*readlink) (struct dentry *, char __user *, int);
```

```
    int (*create) (struct inode *, struct dentry *, umode_t, bool);
```

```
    int (*link) (struct dentry *, struct inode *, struct dentry *);
```

```
    int (*unlink) (struct inode *, struct dentry *);
```

```
    int (*symlink) (struct inode *, struct dentry *, const char *);
```

```
    int (*mkdir) (struct inode *, struct dentry *, umode_t);
```

```
    int (*rmdir) (struct inode *, struct dentry *);
```

```
    int (*mknod) (struct inode *, struct dentry *, umode_t, dev_t);
```

```
    int (*rename) (struct inode *, struct dentry *,
                   struct inode *, struct dentry *, unsigned int);
```

```
    int (*setattr) (struct dentry *, struct iattr *);
```

```
    int (*getattr) (const struct path *, struct kstat *, u32, unsigned int)
```

```
    ssize_t (*listxattr) (struct dentry *, char *, size_t);
```

```
    int (*fiemap)(struct inode *, struct fiemap_extent_info *, u64 start,
                  u64 len);
```

```
    int (*update_time)(struct inode *, struct timespec *, int);
```





```
89 struct dentry {
90     /* RCU lookup touched fields */
91     unsigned int d_flags;           /* protected by d_lock */
92     seqcount_t d_seq;              /* per dentry seqlock */
93     struct hlist_bl_node d_hash;   /* lookup hash list */
94     struct dentry *d_parent;       /* parent directory */
95     struct qstr d_name;
96     struct inode *d_inode;         /* Where the name belongs to - NULL is
97                                     * negative */
98     unsigned char d_iname[DNAME_INLINE_LEN]; /* small names */
99
100    /* Ref lookup also touches following */
101    struct lockref d_lockref;       /* per-dentry lock and refcount */
102    const struct dentry_operations *d_op;
103    struct super_block *d_sb;       /* The root of the dentry tree */
104    unsigned long d_time;           /* used by d_revalidate */
105    void *d_fsdata;                /* fs-specific data */
106
107    union {
108        struct list_head d_lru;     /* LRU list */
109        wait_queue_head_t *d_wait; /* in-lookup ones only */
110    };
111    struct list_head d_child;       /* child of parent list */
112
```





```
67      *      unregistered.
68      */
69
70 int register_filesystem(struct file_system_type * fs)
71 {
72     int res = 0;
73     struct file_system_type ** p;
74
75     BUG_ON(strchr(fs->name, '.'));
76     if (fs->next)
77         return -EBUSY;
78     write_lock(&file_systems_lock);
79     p = find_filesystem(fs->name, strlen(fs->name));
80     if (*p)
81         res = -EBUSY;
82     else
83         *p = fs;
84     write_unlock(&file_systems_lock);
85     return res;
86 }
87
88 EXPORT_SYMBOL(register_filesystem);
89
90 /**
```

