

Архитектура операционной системы

Notes on synchronization

Race conditions

Software defect where result is depended on:

- timings
- uncontrollable events (like scheduling)

- Types of RC

- Static
- Dynamic
- Essential

Critical section

command sequence to access shared data;

- Correctness if:
 - mutual exclusion
 - Progress
 - at least one process can enter if empty
 - Bounded waiting
 - No process wait indefinitely

High level primitives

- Mutex
- Semaphore
- Critical section
- Conditional variables



Linux sync primitives

- Low level
 - Memory barrier
 - Atomic operations
 - Interrupt synchronization
 - Spin locks
- High level
 - Completion
 - Mutex
 - Semaphore
 - Futex

Deadlock conditions

- Mutual exclusion
- Resource waiting
- No resource relocations
- Circle waiting

Deadlock Prevention

- Eliminate one of 4 conditions
 - Mutex
 - Hold and wait
 - May lead to low resource utilization.
 - Starvation is a problem.
 - If it needs additional resources, it releases all of the currently held resources and then requests all of those it needs
 - No Preemption
 - Circular wait



Approaches

- Lock reordering
- Lock manager
- Waiting graph analysis
- Lock-free algorithms

Classical Problems

- Sleeping barber
- Dining Philosophers
- Producer / Consumer

Sleeping barber


The task is modeling queuing system.

- Solution – mutex
- Discussion
 - more than 2 barbers?
- timings



Dining Philosophers

- Problems
 - Starvation
 - Deadlocks & Livelocks
- Solutions
 - Waiter
 - Resource hierarchy
 - Monitors



Producer / consumer problems

The task is to provide synchronized access to shared memory for writers and readers

- Solutions
 - Reader priority
 - Writer priority
 - Fair play

Problems with locking

- Deadlock, livelock
- Priority Inversion
- Convoying

Practice

select

```
int select(int nfds, fd_set *readfds, fd_set *writelfds,  
          fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);  
int  FD_ISSET(int fd, fd_set *set);  
void FD_SET(int fd, fd_set *set);  
void FD_ZERO(fd_set *set);
```

select ex.

```
int
main(void)
{
    fd_set rfd;
    struct timeval tv;
    int retval;

    /* Watch stdin (fd 0) to see when it has input. */
    FD_ZERO(&rfd);
    FD_SET(0, &rfd);

    /* Wait up to five seconds. */
    tv.tv_sec = 5;
    tv.tv_usec = 0;

    retval = select(1, &rfd, NULL, NULL, &tv);
    /* Don't rely on the value of tv now! */

    if (retval == -1)
        perror("select()");
    else if (retval)
        printf("Data is available now.\n");
        /* FD_ISSET(0, &rfd) will be true. */
    else
        printf("No data within five seconds.\n");

    exit(EXIT_SUCCESS);
}
```

Инициализация

Ожидание

Проверка

Домашнее задание

- Курс: <https://stepik.org/course/1780>
 - 4. Средства синхронизации потоков
- Контейнерная виртуализация в Linux